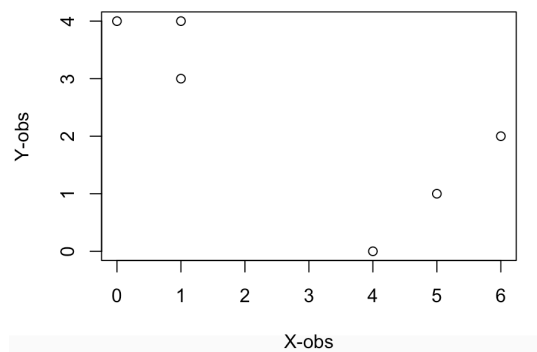


## Performing k-Means by Hand

1. To plot the data, I ran the following code:

```
x <- cbind(c(1, 1, 0, 5, 6, 4), c(4, 3, 4, 1, 2, 0))  
plot(x, xlab="X-obs", ylab="Y-obs")
```

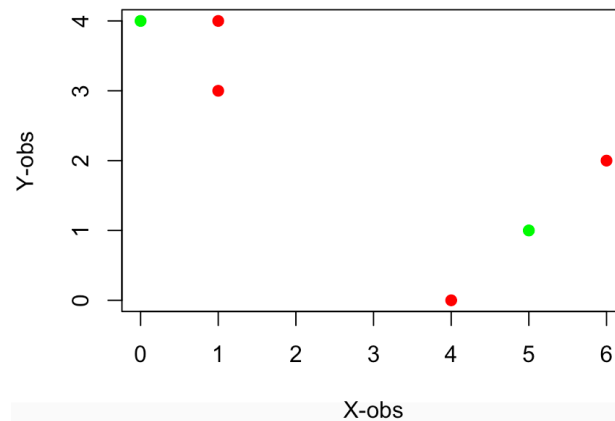
And generated the following graph:



2. For this, I ran the following code:

```
set.seed(17)  
clusters <- sample(seq(0,1), size=6, replace=TRUE)  
plot(x[,1], x[,2], col=ifelse(clusters != 0, 'red', 'green'), pch = 19, xlab="X-obs", ylab="Y-obs")
```

And got the following plot:



3. The centroid for the red cluster is (3.00,2.25) and for the green cluster is (2.5, 2.5). To get these values, I ran the following code:

```
x_green <- mean(x[,1][clusters==0])
y_green <- mean(x[,2][clusters==0])
c(x_green, y_green)
x_red <- mean(x[,1][clusters==1])
y_red <- mean(x[,2][clusters==1])
c(x_red, y_red)
```

4. To determine the distances, I ran the following code:

```
dist_centroid <- function(x, y) {
  dist <- sqrt(sum((x - y) ^2))
  return (dist)
}

centroid_green <- c(x_green,y_green)
centroid_red <- c(x_red,y_red)
cluster_assignment <- c()

for(i in 1:6){
  cluster_assignment[i] <-
    if (dist_centroid(x[i,],centroid_green) <= dist_centroid(x[i,],centroid_red)) 0 else 1
}

cbind(cluster_assignment, x)
```

These are the assignments:

```
      cluster_assignment
[1,]                0 1 4
[2,]                0 1 3
[3,]                0 0 4
[4,]                1 5 1
[5,]                1 6 2
[6,]                1 4 0
```

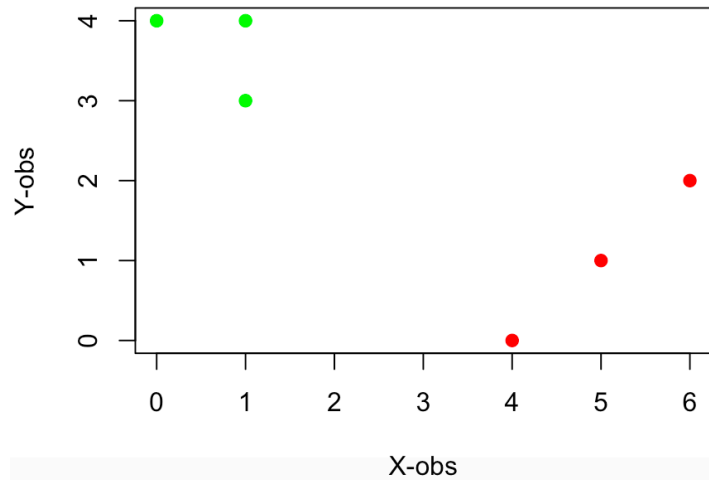
Two of the observations got flipped from red to green, and one from green to red.

5. After running the code from 3 and 4 again, we see that the assignments didn't change. The results are the same as above:

```
      cluster_assignment
[1,]                0 1 4
[2,]                0 1 3
[3,]                0 0 4
[4,]                1 5 1
[5,]                1 6 2
[6,]                1 4 0
```

6. I ran this code to get the plot:

```
plot(x[,1],  
     x[,2],  
     col=ifelse(cluster_assignment != 0, 'red', 'green'),  
     pch = 19,  
     xlab="X-obs",  
     ylab="Y-obs")
```



## Clustering State Legislative Professionalism

1. Loading the data:

```
dat <- load('legprof-components.v1.0.RData')  
dat <- x
```

2. To munge the data, I did the following:

```
dat_a <- select(dat, "t_slength", "slength", "salary_real", "expend", "year", "state")
```

```
dat_b <- subset(dat_a, year==2009 | year==2010)
```

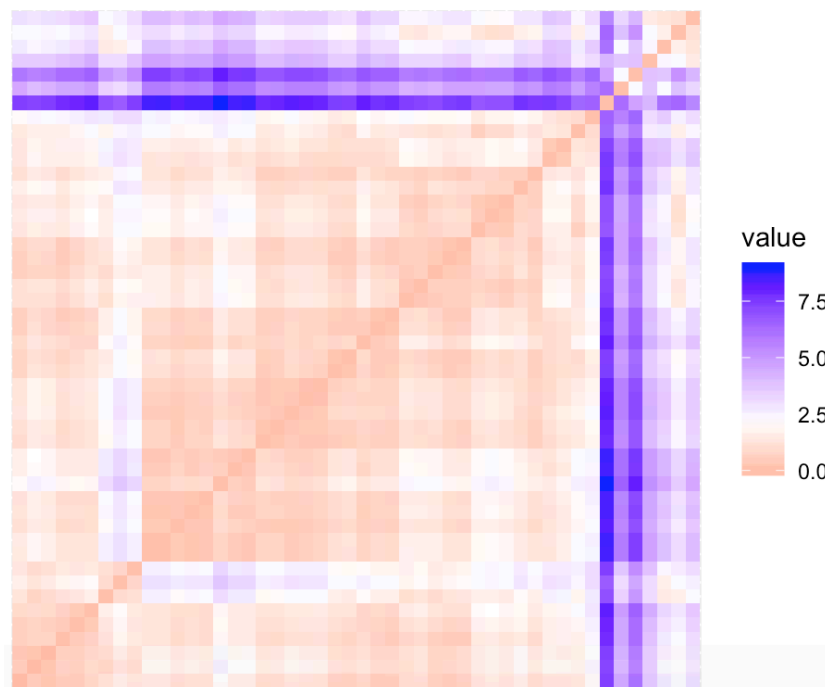
```
dat_c <- na.omit(dat_b)  
state_names = dat_c$state
```

```
dat_d <- select(dat_c, "t_slength", "slength", "salary_real", "expend")  
dat_d <- scale(dat_d)
```

3. To assess clusterability, I used an ODI plot. We can see many adjacent red and white squares in the plot. This, along with our Hopkins stat being close to 1 (.824), we can see that there is some non-random structure and similarity among the data.

```
library(factoextra)
get_clust_tendency(dat_d,n=47)

> get_clust_tendency(dat_d,n=47)
$hopkins_stat
[1] 0.8243719
```

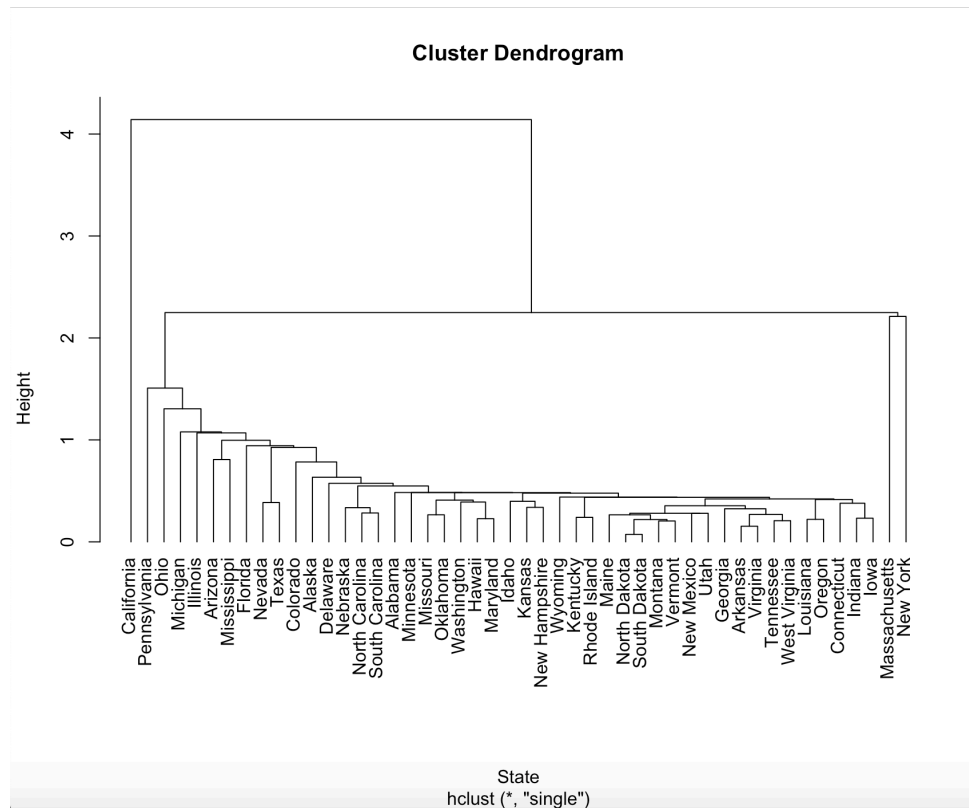


4. To run an agglomerative hierarchical clustering algorithm, I ran the following code:

```
library(tidyverse)
library(skimm)
library(dendextend)

hc_single <- hclust(dist(dat_d),
                    method = "single"); plot(hc_single, hang = -1,xlab="State")
.
```

My results were as follows:



I used the single linkage method. Here, we can see that there are three states that are very dissimilar to the others: California, Massachusetts, and New York. We can see that, not surprisingly, North and South Carolina and North and South Dakota are very similar. Indiana and Iowa are another similar pair, which makes sense due to the geographic similarity and location. An unusual, or unexpected, similarity is between Hawaii and Maryland. The only patterns I see are with location. I can't say why Hawaii and Maryland are similar, but it makes sense that North and South Dakota are similar.

5. To run a k-means fit, I ran the following code:

```
kmeans <- kmeans(dat_d,
                  centers = 2,
                  nstart = 30)

t <- as.table(kmeans$cluster)
t <- data.frame(t)
clusterk <- t[t$Freq == "2",]
clusterk

str(kmeans)

kmeans$cluster
kmeans$centers
kmeans$size
```

The results were:

```

> clusterk
      Var1 Freq
5    California    2
21 Massachusetts    2
22     Michigan    2
31     New York    2
34         Ohio    2
37 Pennsylvania    2

      t_slength    slength salary_real    expend
1 -0.2868507 -0.2949065    -0.29189 -0.2092542
2  2.0079549  2.0643454     2.04323  1.4647791
> kmeans$size
[1] 42  6

```

What we see here is that the two clusters are composed of 42 states, and 6 states. The 6 states in cluster 2 are shown above. Looking at the breakdown of the clusters, we can see that the states in cluster 2 have much higher values in each variable of interest.

6. To run the GMM, I ran the following code:

```

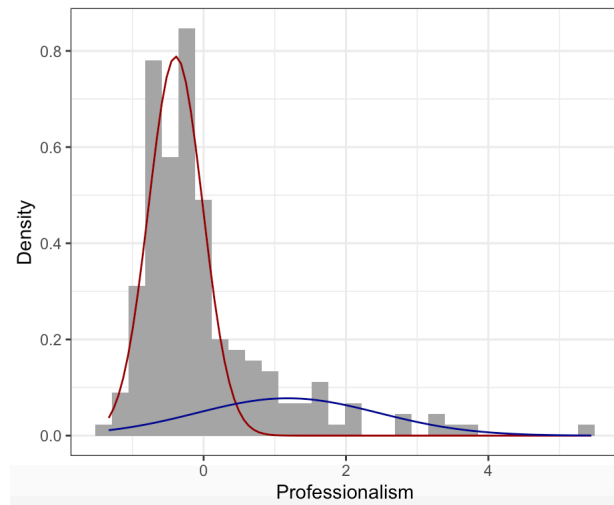
library(mixtools)
library(plotGMM)
set.seed(1234)

gmm1 <- normalmixEM(dat_d, k = 2)

ggplot(data.frame(x = gmm1$x)) +
  geom_histogram(aes(x, ..density..), fill = "darkgray") +
  stat_function(geom = "line", fun = plot_mix_comps,
    args = list(gmm1$mu[1], gmm1$sigma[1], lam = gmm1$lambda[1],
    colour = "darkred") +
  stat_function(geom = "line", fun = plot_mix_comps,
    args = list(gmm1$mu[2], gmm1$sigma[2], lam = gmm1$lambda[2],
    colour = "darkblue") +
  xlab("Professionalism") +
  ylab("Density") +
  theme_bw()

```

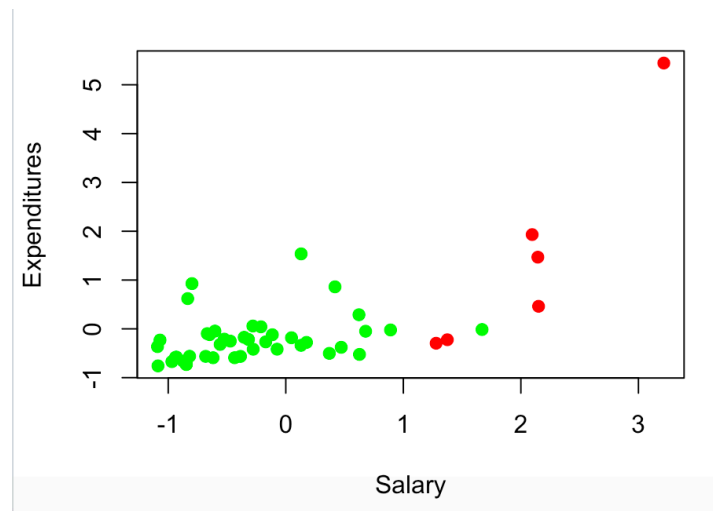
This was the resulting plot:



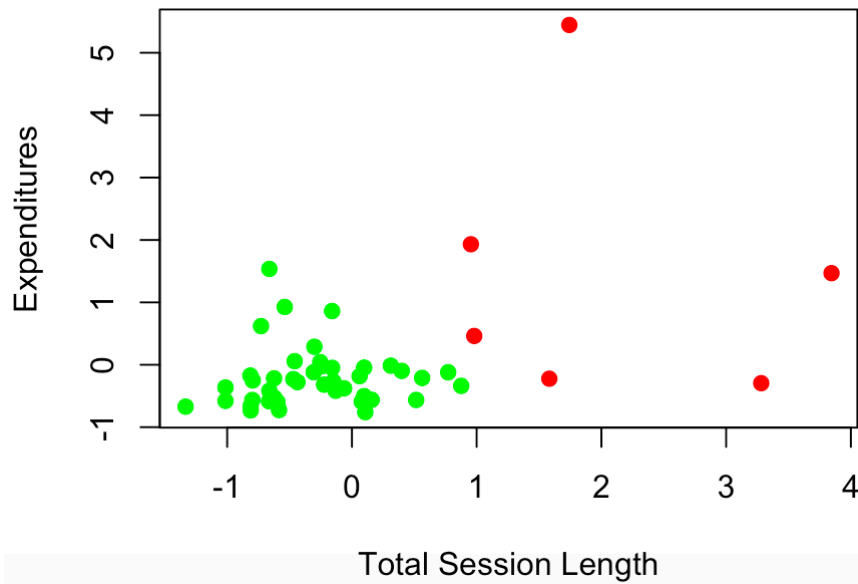
From this, we can see that there are again two distinct groups. One tends to more professionalism, but with much lower density (or frequency), than the other group which is much less professional with higher density. This reflects what we found in the k-means clustering.

- First, we can see that on the dendrogram above, the 6 states that are most dissimilar are the same as we saw in the k-means clustering algorithm. To further explore this, I plotted the values of different variables and colored them based on assignment.

```
plot(dat_d[,3],
     dat_d[,4],
     col=ifelse(assignments != 0, 'red', 'green'),
     pch = 19,
     xlab="Salary",
     ylab="Expenditures")
```



```
plot(dat_d[,2],
     dat_d[,4],
     col=ifelse(assignments != 0, 'red', 'green'),
     pch = 19,
     xlab="Total Session Length",
     ylab="Expenditures")
```



From these graphs, we see that there are two fairly distinct groups of states, and that the clustering of these groups makes sense. If we compare key values from the k-means and GMM models, we see this:

```
> kmeans$centers
  t_slength  slength salary_real  expend
1 -0.2868507 -0.2949065   -0.29189 -0.2092542
2  2.0079549  2.0643454    2.04323  1.4647791
> gmm1$mu
[1] -0.3919019  1.1807282
```

We see here that the centers for each cluster in both models is somewhat similar in magnitude and sign.

8. To run this validation check, I used the clValid library:

```
library(clValid)
library(mclust)
validation_check <- clValid(dat_d,
                             c(2),
                             clMethods = c("hierarchical", "kmeans", "model"),
                             validation = c("internal")); summary(validation_check)
```



The outputs were:

```
Clustering Methods:
hierarchical kmeans model

Cluster sizes:
2

Validation Measures:
2

hierarchical Connectivity 6.0869
                  Dunn    0.3598
                  Silhouette 0.6920
kmeans             Connectivity 8.5683
                  Dunn    0.1726
                  Silhouette 0.6390
model              Connectivity 18.7095
                  Dunn    0.0833
                  Silhouette 0.4230

Optimal Scores:

          Score Method    Clusters
Connectivity 6.0869 hierarchical 2
Dunn          0.3598 hierarchical 2
Silhouette    0.6920 hierarchical 2
```

What we see here is that The silhouette scores for each model are fairly similar, with values of .69, .63, and .42. What this tells us is that first, the hierarchical model performed the best, with k-means coming in a close second, and then the GMM model following in last. If we run this with different cluster sizes, we get the following:

```
validation_check <- clValid(dat_d,
                             c(2:10),
                             clMethods = c("hierarchical", "kmeans", "model"),
                             validation = c("internal")); summary(validation_check)
```

Clustering Methods:  
hierarchical kmeans model

Cluster sizes:  
2 3 4 5 6 7 8 9 10

Validation Measures:		2	3	4	5	6	7	8	9	10
hierarchical	Connectivity	6.0869	6.9536	13.1345	15.1345	20.7563	22.9230	28.1726	30.1171	40.5512
	Dunn	0.3598	0.4340	0.2902	0.2902	0.2836	0.2836	0.2451	0.2451	0.1930
	Silhouette	0.6920	0.6619	0.5199	0.4989	0.3776	0.3658	0.2921	0.2831	0.2624
kmeans	Connectivity	8.5683	11.0183	18.1651	20.1651	23.6810	25.8476	36.4726	44.8750	45.4024
	Dunn	0.1726	0.2597	0.2456	0.2456	0.1214	0.1214	0.1871	0.1846	0.2515
	Silhouette	0.6390	0.6054	0.4824	0.4611	0.3328	0.3210	0.3169	0.2854	0.3249
model	Connectivity	18.7095	23.7964	33.3683	60.1651	69.0651	54.4433	51.8206	63.7619	62.1766
	Dunn	0.0833	0.0855	0.0554	0.0280	0.0391	0.0532	0.0935	0.0879	0.0928
	Silhouette	0.4230	0.3854	0.2157	0.0962	0.0473	0.1822	0.2957	0.2091	0.2132

Optimal Scores:

	Score	Method	Clusters
Connectivity	6.0869	hierarchical	2
Dunn	0.4340	hierarchical	3
Silhouette	0.6920	hierarchical	2

As we can see, the silhouette scores decrease for each model after 2 clusters, which tells us that k=2 is the optimal number of clusters for this data.

- From all of this analysis, what we see is that using the silhouette scores, all of our models perform best on k=2 clusters. However, if we look at the Dunn score, our hierarchical model performs best at k=3 clusters. This tells us that our model may not be perfect at k=2. However, with the silhouette score, the hierarchical model score is fairly close to 1 (.692), which indicates a better fit. It seems that our hierarchical model is best at a specification of k=2.

We may select a “sub-optimal” clustering method due to the way a clustering algorithm works. If we select k=n(obs), then we’d have a perfectly fit model. However, the variance would be very high and the model wouldn’t actually tell us anything. If we can use a smaller amount of clusters, resulting in an imperfect model, it still may be much more informative than a “perfectly fit” model as described above. There is a tradeoff in choosing parameters, and often, we must choose a “sub-optimal” clustering method in order to have results that are informative.